





Reordering Effects

This lesson discusses the compiler, runtime or hardware optimizations that can cause reordering of program instructions

Take a look at the following program and try to come up with all the possible outcomes for the variables **ping** and **pong**.

```
1  class Demons
2      public s
3      (new R
4      }
5  }
6
7  class Reorde
8
9      private
10     private
11     private
12     private
13
14     public v
15
16         Thre
17
18
19
20
21
22     });
23
24         Thre
25
26
27
28
```



Effect of reordering on program output

Most folks would come up with the following possible outcomes:

- 1 and 1
- 1 and 0
- 0 and 1

However, it might surprise many but the program **can very well print 0 and 0!** How is that even possible? Think from the point of view of a compiler, it sees the following instructions for **thread t1's run()** method:

```
bar = 1;  
pong = foo;
```

The **compiler doesn't know that the variable bar is being used by another thread so it may take the liberty to *reorder* the statements** like so:

```
pong = foo;  
bar = 1;
```

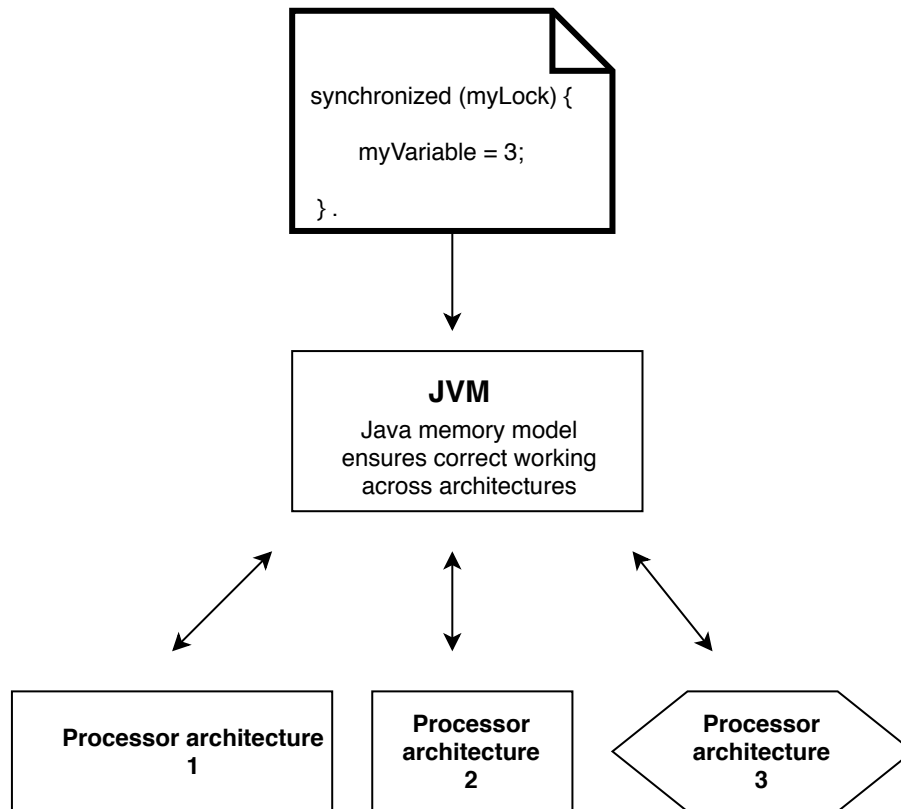
The two statements don't have a dependence on each other in the sense that they are working off of completely different variables. For performance reasons, the compiler may decide to switch their ordering. Other forces are also at play, for instance, the **value of one of the variables may get flushed out to the main memory from the processor cache but not for the other variable.**



Note that with the reordering of the statements the JVM still is able to honor the *within-thread as-if-serial* semantics and is completely justified to move the statements around. Such performance and optimization tricks by the compiler, runtime or hardware catch unsuspecting developers off-guard and lead to bugs which are very hard to reproduce in production.

Platform

Java is touts the famous *code once, run anywhere* mantra as one of its strengths. However, this isn't possible without Java shielding us from the vagrancies of the multitude of memory architectures that exist in the wild. For instance, the frequency of reconciling a processor's cache with the main memory depends on the processor architecture. A processor may relax its memory coherence guarantees in favor of better performance. The architecture's memory model specifies the guarantees a program can expect from the memory model. It will also specify instructions required to get additional memory coordination guarantees when data is being shared among threads. These instructions are usually called *memory fences or barriers* but the Java developer can rely on the JVM to interface with the underlying platform's memory model through its own memory model called **JMM** (Java Memory Model) and insert these platform memory specific instructions appropriately. Conversely, the JVM relies on the developer to identify when data is shared through the use of proper synchronization.



Interviewing soon? We've partnered with Hired so that companies apply to you
[utm_source=educative&utm_medium=lesson&utm_location=CA&utm_campaign=educative](https://www.hired.com/?utm_source=educative&utm_medium=lesson&utm_location=CA&utm_campaign=educative)

[← Back](#)[Next →](#)

Java Memory Model

The happens-before Relationship



Mark as Completed

Report an
Issue

Ask a Question

(https://discuss.educative.io/tag/reordering-effects__java-memory-model__java-multithreading-for-senior-engineering-interviews)

